

# As Much Resilience As You Want: Application-tailored Resilience in Legion

**Abstract**—Resilient is an important aspect of scaling applications on large clusters. As we approach parallelism profiles of several millions and long running applications, we need to ensure that ineffect “we reach the end”. Defining the policy interaction with the programming model features necessitates that we revisit the memory consistency model. Recoverability, a critical step in resilience, opens the door to optimizations such as speculation. We also evaluate this. ...

**Index Terms**—resilience, legion

## I. INTRODUCTION

what is the need for resilience

what are the challenges we face with resilience in a deferred execution model, with all the decisions taken at runtime via a mapper ?

what are the advantages that we derive from the same ?  
The needle in the resilience spectrum can be much finer than ABFT, sub-tree DAG based, snapshot based, relocatable-finish based, strategies. Dynamic and deferred being the key.

what are the advantages that we derive from having multiple wavefront ?  
how do we ensure minimal overhead ?

Is it really a resilience framework, or a glorified garbage collector ?

how is it different from x10, parsec, charm++  
- they also allow tasks to be marked as resilience  
- x10 allows finish blocks, exception semantics.  
- what are the exception semantics that we are providing

what about regent ? where do we stand there ?

what about local vs global recovery  
- can we recover from a node failure  
- can we recover from an exception, what are the semantics that we provide here  
- can we recover from ECC errors,  
- can we fix these errors ?  
- can we recover from I/O errors ?  
- can we recover from non-SingleTasks, what is recovery for index space tasks, must epoch tasks

what is the memory consistency model ? what is the state of the regions, the tasks that are physically dependent but are

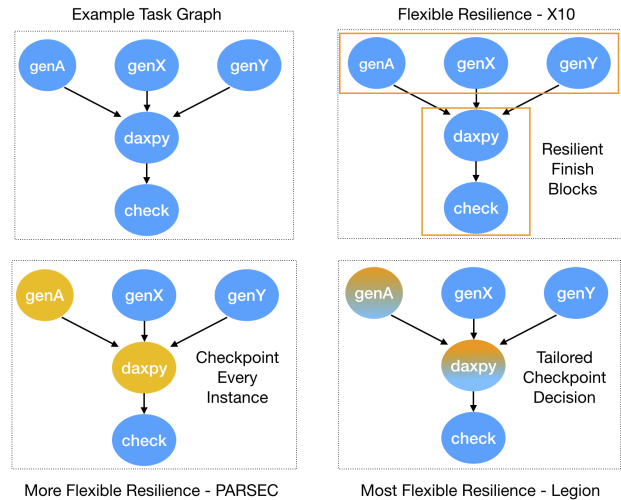


Fig. 1. A spectrum of resilience policies supported by different state-of-the-art parallel runtimes.

not logically dependent ?

experiments: circuit, miniAero, Soleil-X, stencil, (S3D ?)

## II. OVERVIEW OF RESILIENCE IN LEGION

### III. INTERACTION OF RESILIENCE POLICY WITH LEGION FEATURES

which are the legion features of importance ?

put a figure of interaction.

what is the memory consistency angle here ?

what does it mean to advance the commit wavefront

**Resilience is a tangling of the lifetime of a task and a region snapshot**

1) when can we advance a commit wavefront ? 2) what's the lifetime of a region-instance snapshot ?

on this front, we see it as a two-step process a) define a consistent cut of tasks problem, b) commit any task strictly behind this wavefront c) garbage collect any snapshot that serves as input to any task that is already committed

consistent cut of tasks that can be part of the commit wavefront: a set of tasks whose inputs are need\_preserve'd, or they are strictly post-dominated by tasks whose inputs are need\_preserved.

3) what about copy/index/tasks ? copy local to local follows the above semantics copy local to remote get committed immediately after successful execution. index launch tasks,

actually feel need not be in the task graph, unless virtual mapping is used. they can be garbage collected immediately after all child tasks are included in the dependence analysis wavefront - I am thinking we will never have a case where are relaunching the index launch, since the child tasks either are part of the committed wavefront or are not (pending a discussion on phase barriers)

**What to do about must\_epoch tasks with phase\_barrier inside them ?**

answer: restartable phase\_barrier with generation commit callback

#### IV. IMPLEMENTATION OF RESILIENCE

##### A. *need\_preserve* Implementation

tagging region instances tagging tasks

-the mapper marks an instance as persistent, i.e., vector<vector<bool>> persistent -in finalize\_map\_task, the single-Task sees this and notes it down in the Individual\_task's persistent\_tasks list -when trigger\_complete is called on the task,

1) inside line 5560, invalidate\_region\_tree\_contexts, inside which we have runtime->forest->invalidate\_versions, we do not do on region[idx].region. we also do not do the instance\_top\_views

2) we retain the mark on the task as allowed\_for\_gc.

3) we go to the incoming of this task, and just like verified\_regions[true], we mark outgoing\_edge\_dominated[true] if all the outgoing of a task is marked as true, then we change allowed\_for\_gc = true.

Discussion Points for today ————— 1) allow persistence on a subset of mapped instances in map\_task ? Pro: flexibility, Con: if a checkpoint is to be considered useful for a restart, not having the full set of inputs checkpointed seems contradictory. 2) verify\_regions tracks op dependencies after physical dependence analysis, correct ? 3) a discussion on persistence inside map\_task call

discussed design:

1) build a function set\_hardened\_instance(instance,task) along the lines of the set\_gc\_priority(instance, never, task) inside the mapper call 2) inside that mark a task's incoming edges as saying that it leads to a hardened instance(task) 3) the garbage collector will basically collect a task whose outgoing edges are all marked as verified/hardened. 4) if a task is gc'ed, then it marks all its incoming edges as leads to a hardened instance. 5) steps 1-4 will be based on set\_garbage\_collection\_priority and how verified\_regions are set. We will be adding a new list to each task, similar to verified\_regions, that will represent edges\_ending\_in\_hardened\_regions.

quash\_operation: 847 line in realm/proc\_impl.cc tells us that we should set the poisoned to true and ensure that the start\_event.has\_triggered\_faultaware(poisoned) returns true

what about two sibling tasks that map the same instance, and one hardens it, the other does not - alex raised this, we discussed that trigger recover would check - but the way we are doing it now, we are doing unverified\_regions.erase based

on even on guy below us who will call harden on a region instance - so the parent would have gone away, but the child B who did not call harden on the physical region should be able to restart since the instance will be there how to show this is working, if the application is too fast, then delete meta task on the logical region is called which deletes the region irrespective of whether its GC\_NEVER\_PRIORITY

REACTIVE faulty task catch it poison propagation identify recovery point trigger\_recover

SEMANTICS regions and instances and consistency

PROACTIVE mark instances as harden

DIFFERENCE WITH PARSEC: label a task with snapshot, i.e., all or none policy whereas we allow selective instances to be marked as harden. - what is an example app that is benefited by this ?

1) *Interaction of need\_preserve with the commit wavefront:*

2) *Obtaining dependence graph in the mapper before calling need\_preserve:*

##### B. *ProfilingResponse* callback used

while using the profiling measurement reporting infrastructure as-is causes overhead, this is because of the invoking of a mapper profiling response function. We do not need that for resilience, so the resilience callback would short circuit it.

task launch -> profiling reported event is there -> when it is triggered -> singletask::profiling\_response is invoked -> handle things — in here, there is no need to call the mapper side yet. So, avoid one overhead there, maybe this was never called and so , this is not an optimization. ok, back to square one.

#### V. RESILIENCE APPLICATION: SPECULATION

there are three different wavefronts in Legion, we can speculate on any of them.

From a speculation perspective, are they different ? Are we novel, since we have these three different wavefronts ? Can we navigate through this, like the blanks

1) execution wavefront what does it mean to speculate here, does the other steps have to be complete before we do this.

2) mapping wavefront

3) dependence analysis wavefront

– see mike's 6 wavefront answer.

There are three different wavefronts, there could

**the 6 wavefronts, where does speculation plays a role**

mike - speculation is more about tracking the resolution wavefront while resilience is more about tracking the commit wavefront, but the when things go bad, then i think the machinery to restart the mapping and execution wavefronts should be the same

#### VI. EXPERIMENTS

##### A. *Index Tasks*

a) *Growth of memory as execution proceeds:*




Fig. 2. Total time taken by 02\_index\_tasks. Time in Seconds, 1000 tasks/index launch




Fig. 3. Total time taken by 02\_index\_tasks. Time in Seconds, 100 tasks/index launch

b) Performance overhead as execution proceeds:

NumIter	Native	Resl	lg:resl
128	23	16	21
256	20	22	28
512			

NumIter	Native	lg:resl	Resl
100	7.2	6.6	23.6
200	13.8	13.3	46.3
400	26.1	26.5	94.4
800	55	53	186.8
1000	65	66	239.9

NumIter	Native	Resilience
100	1.12	2.43
200	1.56	4.39
400	3.00	9.45
800	5.58	19.09
1000	7.34	23.2

NumIter	Native	lg:resl	Resl
100	18	140	107
200	22	263	176
400	24	509	245
800	26	1002	428
1000	26	1248	518




Fig. 4. Total Memory footprint by 02\_index\_tasks. Memory in MB, 1000 tasks/index launch.

B. Stencil

C. local recover vs global recovery

D. compute/comm vs no-failure/single failure/multi-failure

E. S3D, Pennant, Stencil, Circuit

F. Some Interesting Task Graphs for Recovery

G. bigger examples

pennant, stencil, miniAero (better understood) /Circuit, - limit the failure to the before

## VII. ADAPTIVE RESILIENCE

### VIII. RESILIENCE POLICY IMPLEMENTED VIA MAPPER: EXAMPLE 1 GENERIC

### IX. RESILIENCE POLICY IMPLEMENTED VIA MAPPER: EXAMPLE 2 UT AUSTIN

## X. CONCLUSION

## ACKNOWLEDGMENT

## REFERENCES

- [1] S. Treichler, M. Bauer, and A. Aiken. Language support for dynamic, hierarchical data partitioning. In Object Oriented Programming, Systems, Languages, and Applications (OOPSLA), 2013.